

Twelve binary items in a fixed order: the statistical gate first, the economic gate second.

Stage 1 · Statistical verifiable from the manuscript alone

- ☐ **V1. Report prediction class distribution**
Do you report the class distribution of your model's predictions?
- ☐ **V2. Test with/without class balancing**
Do you test both with and without class balancing?
- ☐ **V3. Use temporal splitting only**
Do you split strictly by time, with purge and embargo gaps?
- ☐ **V4. Apply CPCV with PBO**
Do you apply CPCV and report the probability of backtest overfitting?
- ☐ **V5. Report Var(SR_IS)**
Do you report the variance of in-sample Sharpe across configurations?
- ☐ **V6. Include permutation tests (≥ 100)**
Do you run a permutation test with at least 100 iterations?

Stage 2 · Economic required before anything is traded

- ☐ **V7. Report net performance with costs**
Do you report performance net of transaction costs?
- ☐ **V8. Cost sensitivity analysis**
Do you show how conclusions hold across a range of cost assumptions?
- ☐ **V9. Compare against simple baselines**
Do you compare against simple active baselines, not only buy-and-hold?
- ☐ **V10. Evaluate in bear markets and regime transitions**
Do you evaluate in bear markets and regime transitions?
- ☐ **V11. Report trade frequency**
Do you report trade frequency?
- ☐ **V12. Provide code for reproducibility**
Is your code available so the result can be reproduced?

How to score

One point per item, satisfied or not — no partial credit. Score Stage 1 and Stage 2 separately, then total out of 12.

What the total means

The median paper in an audit of 74 empirical crypto ML studies scores 2.5 of 12; none of the 74 applied CPCV. A study that clears Stage 1 but fails Stage 2 is the statistical-economic disconnect: a significant signal with no exploitable alpha.

What it is not

Self-scoring is not a certification. It records your own reading of your own artifacts. We scored our own camera-ready 9 of 12 and disclosed three partials.